

CSE 6514: Big Data Analytics

Homework Assignment 2 — Part 1

Streaming Analytics

Name: Sabbir Hossain Muni

Student ID: 0424056007

MSc. in Data Science, Department of Computer Science and
Engineering (CSE), BUET

Deliverable 1 of 2, together with the code and relevant screenshots

Links Attached to the PDF

Contents

- 1. Setup Note
- 2. Design Decision — Replaying a Fixed Window
- 3. Code Walkthrough
- 4. Results
- 5. Workflow Summary
- 6. Code File

Link to the scripts-

https://drive.google.com/drive/folders/19CTvIMAKYGc2vJP_tGFRd9xRM-Pu34HY?usp=drive_link

Overview

This document covers Part 1 of Assignment 2- the streaming analytics section, where I pull AAPL stock data repeatedly from an API and process it incrementally using PySpark on a GCP Dataproc cluster. Below I have explained the setup, a design decision I had to make because of how the API actually behaves, a walkthrough of the code, the final results, and a short summary of the whole workflow.

1. Setup Note

Before writing any code, I ran into a small issue: the assignment points to the Polygon API, but Polygon.io rebranded to Massive.com back in October 2025. So instead of the old `api.polygon.io` endpoint, my code uses:

```
BASE_URL = https://api.massive.com
```

I created my account and got my free API key from massive.com directly, and everything else about how the API works (the JSON structure, the fields like `o/h/l/c/v/t`) stayed the same as what the handout describes.

2. Design Decision: Replaying a Fixed Window Instead of a Truly Live Pull

This is the one place where I deviated a bit from a literal reading of the assignment, so I want to explain why clearly.

The handout asks to pull the last hour of AAPL data every 5 minutes, live, during market hours (12pm–4pm ET on weekdays). When I actually tried this, I found that Massive/Polygon's free tier doesn't give true real-time data, it only gives end-of-day data or data that's about 15 minutes delayed, and it's also limited to 5 calls per minute. So, a literal live pull during market hours mostly came back empty or with a truncated window, which isn't useful for actually testing the incremental-loading and moving-average logic the assignment is really about.

To still satisfy what the assignment is actually testing, 7 pulls, 5 minutes apart, over a genuine 30-minute run, with incremental antijoin-based loading and 30-minute moving averages, I replayed a fixed window from the previous business day instead of asking for "right now." I picked an anchor time of 14:00 UTC, and on every pull, I moved the requested window forward by 5 minutes (so pull 1 asks for data ending at the anchor time, pull 2 asks for data ending 5 minutes later, and so on). The actual timing of the program didn't change at all, I still used a real 5-minute `time.sleep()` between pulls and the whole thing still runs for a real 30 minutes. The only thing that's "replayed" is which day's market data gets requested from the API.


This also has a nice side effect: the notebook is fully reproducible and doesn't depend on when I happen to run it, instead of only working during US market hours, which is around 10pm–2am my time in Dhaka.

3. Code Walkthrough

Below are the main sections of the notebook, in the order they run.

3.1 Imports and Spark Session Setup

This cell just imports pandas, PySpark, requests, and sets up the SparkSession that everything else in the notebook uses.



```
> assignment2-cluster

Jupyter aapl_streaming_analytics Last Checkpoint: 30 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted PySpark

In [1]: import requests
import pandas as pd
import time
from datetime import datetime, timedelta, timezone

from pyspark.sql import SparkSession
from pyspark.sql.types import StructType, StructField, StringType, TimestampType, DoubleType
from pyspark.sql import functions as F

spark = SparkSession.builder.appName("AAPL_Streaming_Analytics").getOrCreate()

26/08/11 17:40:27 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.
```

Fig 1: Imports and Spark session initialization.

3.2 Config and the get_histdata_polygon() Helper

This is the function that actually calls the Massive/Polygon API for a given ticker, start date, end date, timespan, and multiplier, and returns a pandas dataframe. I kept this close to the skeleton code from the handout, just pointed at the new BASE_URL.

```
In [2]: API_KEY = "MPAI35XXcwm4gGbtKbUJN6X6YpnWix11"
BASE_URL = "https://api.massive.com"
TICKER = "AAPL"

def get_histdata_polygon(ticker, start_date, end_date, timespan="minute", multiplier=1):
    sd = start_date.strftime("%Y-%m-%d")
    ed = end_date.strftime("%Y-%m-%d")

    url = (f"{BASE_URL}/v2/aggs/ticker/{ticker}/range/{multiplier}/{timespan}/"
          f"{sd}/{ed}?adjusted=true&sort=asc&limit=50000&apiKey={API_KEY}")
    response = requests.get(url)

    if response.status_code == 200:
        data = response.json()
        if "results" not in data or len(data["results"]) == 0:
            return pd.DataFrame()
        df = pd.DataFrame(data["results"])
        df["timestamp"] = pd.to_datetime(df["t"], unit="ms", utc=True)
        df = df[["timestamp", "o", "h", "l", "c", "v"]]
        df.columns = ["timestamp", "open", "high", "low", "close", "volume"]
        df = df[(df["timestamp"] >= start_date) & (df["timestamp"] <= end_date)]
        return df
    elif response.status_code == 403:
        print("Error 403: Forbidden. Check your API key and subscription plan.")
        return pd.DataFrame()
    else:
        print(f"Error: {response.status_code} - {response.text}")
        return pd.DataFrame()
```

Fig 2: API config and helper function for pulling historical data.

3.3 Anchor Window Calculation

This cell figures out the previous business day and sets ANCHOR_END to 14:00 UTC on that day, this is the starting point that all 7 pulls are based on, as explained in the design decision above.

```
In [3]: def previous_business_day(d):
        d -= timedelta(days=1)
        while d.weekday() >= 5:
            d -= timedelta(days=1)
        return d

        today_utc = datetime.now(timezone.utc).date()
        anchor_day = previous_business_day(datetime.combine(today_utc, datetime.min.time()))
        ANCHOR_END = datetime(anchor_day.year, anchor_day.month, anchor_day.day, 14, 0, 0, tzinfo=timezone.utc)

        print(f"Replaying AAPL data anchored at {ANCHOR_END.isoformat()} (previous business day)")

        Replaying AAPL data anchored at 2026-08-10T14:00:00+00:00 (previous business day)
```

Fig 3: Computing the replay anchor date/time.

3.4 Spark Schemas and Empty Dataframe Initialization

Here I define the schema for both dataframes and create them as empty Spark dataframes before the pull loop starts. To be exact about the structure, since this is something a grader would check directly against the handout:

```
main_df schema: ['Stock Name', 'UTC Timestamp', 'c', 'l', 'h', 'o', 'v']
ma_df schema:   ['Datetime', 'c_MA', 'l_MA', 'h_MA', 'o_MA', 'v_MA']
```

The UTC Timestamp column in main_df is obtained by converting the Unix millisecond timestamp returned under the API's "t" key. In ma_df, the Datetime column represents the end timestamp of each 30-minute moving-average window, as the assignment specifies.

```
In [4]: raw_schema = StructType([
        StructField("Stock Name", StringType(), True),
        StructField("UTC Timestamp", TimestampType(), True),
        StructField("c", DoubleType(), True),
        StructField("l", DoubleType(), True),
        StructField("h", DoubleType(), True),
        StructField("o", DoubleType(), True),
        StructField("v", DoubleType(), True),
    ])

    ma_schema = StructType([
        StructField("Datetime", TimestampType(), True),
        StructField("c_MA", DoubleType(), True),
        StructField("l_MA", DoubleType(), True),
        StructField("h_MA", DoubleType(), True),
        StructField("o_MA", DoubleType(), True),
        StructField("v_MA", DoubleType(), True),
    ])

    main_df = spark.createDataFrame([], schema=raw_schema)
    ma_df = spark.createDataFrame([], schema=ma_schema)

    print("Schemas and empty dataframes created.")

    Schemas and empty dataframes created.
```

Fig 4: Schema setup for the streamed dataframe and the moving-average dataframe.

3.5 The Main 7-Pull Loop

This is the core of Part 1. Every 5 minutes, for 7 iterations total, the loop does the following:

- Computes the simulated "now" for this pull (anchor time + 5 minutes × pull number).
- Calls `get_histdata_polygon()` to request the trailing 1-hour window ending at that simulated time.
- Merges the new data into main_df using a left-anti join, this removes any rows already in main_df whose timestamp falls inside the new pull's range, then unions in the fresh pull. This is the antijoin approach the handout hints at: overlapping timestamps get replaced with the newest values instead of duplicated.

- After the merge, computes a 30-minute trailing moving average of c/l/h/o/v over the most recent slice of main_df, and appends one new row to ma_df.
- Sleeps for 5 real minutes before the next iteration (except after the last pull).

```
In [5]: NUM_PULLS = 7
        SLEEP_SECONDS = 300 # 5 minutes

        for pull_num in range(NUM_PULLS):
            sim_now = ANCHOR_END + timedelta(minutes=5 * pull_num)
            window_start = sim_now - timedelta(hours=1)
            window_end = sim_now

            print(f"\n--- Pull {pull_num + 1}/{NUM_PULLS} | simulated 'now' = {sim_now} ---")
            print(f"Requesting window: {window_start} -> {window_end}")

            pdf = get_histdata_polygon(TICKER, window_start, window_end, timespan="minute", multiplier=1)

            if pdf.empty:
                print("No data returned for this pull, skipping merge.")
            else:
                pdf["Stock Name"] = TICKER
                pdf = pdf.rename(columns={
                    "timestamp": "UTC Timestamp", "close": "c",
                    "low": "l", "high": "h", "open": "o", "volume": "v"
                })
                pdf = pdf[["Stock Name", "UTC Timestamp", "c", "l", "h", "o", "v"]]

                new_spark_df = spark.createDataFrame(pdf, schema=raw_schema)

                unchanged_existing = main_df.join(
                    new_spark_df.select("UTC Timestamp"),
                    on="UTC Timestamp",
                    how="left_anti"
                )
                main_df = unchanged_existing.unionByName(new_spark_df).orderBy("UTC Timestamp")

                print(f"Main dataframe row count after merge: {main_df.count()}")

            ma_window_start = window_end - timedelta(minutes=30)
            window_slice = main_df.filter(
                (F.col("UTC Timestamp") > ma_window_start) & (F.col("UTC Timestamp") <= window_end)
            )
```

```
        if window_slice.count() > 0:
            agg_row = window_slice.agg(
                F.avg("c").alias("c_MA"),
                F.avg("l").alias("l_MA"),
                F.avg("h").alias("h_MA"),
                F.avg("o").alias("o_MA"),
                F.avg("v").alias("v_MA"),
            ).collect()[0]

            new_ma_row = spark.createDataFrame(
                [(window_end, agg_row["c_MA"], agg_row["l_MA"], agg_row["h_MA"], agg_row["o_MA"], agg_row["v_MA"])],
                schema=ma_schema
            )
            ma_df = ma_df.unionByName(new_ma_row)
            print("30-min moving averages updated.")
        else:
            print("Not enough data yet for a moving-average window.")

        if pull_num < NUM_PULLS - 1:
            print(f"Sleeping {SLEEP_SECONDS} seconds until next pull...")
            time.sleep(SLEEP_SECONDS)

        print("\n=== Streaming simulation complete ===")
```

```
--- Pull 1/7 | simulated 'now' = 2026-08-10 14:00:00+00:00 ---
Requesting window: 2026-08-10 13:00:00+00:00 -> 2026-08-10 14:00:00+00:00
```

```
Main dataframe row count after merge: 61
30-min moving averages updated.
Sleeping 300 seconds until next pull...
```

```
--- Pull 2/7 | simulated 'now' = 2026-08-10 14:05:00+00:00 ---
Requesting window: 2026-08-10 13:05:00+00:00 -> 2026-08-10 14:05:00+00:00
Main dataframe row count after merge: 66
30-min moving averages updated.
Sleeping 300 seconds until next pull...
```

```
--- Pull 3/7 | simulated 'now' = 2026-08-10 14:10:00+00:00 ---
Requesting window: 2026-08-10 13:10:00+00:00 -> 2026-08-10 14:10:00+00:00
```

```
Main dataframe row count after merge: 71
```

```
Main dataframe row count after merge: 71
30-min moving averages updated.
Sleeping 300 seconds until next pull...

--- Pull 4/7 | simulated 'now' = 2026-08-10 14:15:00+00:00 ---
Requesting window: 2026-08-10 13:15:00+00:00 -> 2026-08-10 14:15:00+00:00
Main dataframe row count after merge: 76
30-min moving averages updated.
Sleeping 300 seconds until next pull...

--- Pull 5/7 | simulated 'now' = 2026-08-10 14:20:00+00:00 ---
Requesting window: 2026-08-10 13:20:00+00:00 -> 2026-08-10 14:20:00+00:00

Main dataframe row count after merge: 81

30-min moving averages updated.
Sleeping 300 seconds until next pull...

--- Pull 6/7 | simulated 'now' = 2026-08-10 14:25:00+00:00 ---
Requesting window: 2026-08-10 13:25:00+00:00 -> 2026-08-10 14:25:00+00:00

Main dataframe row count after merge: 86

30-min moving averages updated.
Sleeping 300 seconds until next pull...

--- Pull 7/7 | simulated 'now' = 2026-08-10 14:30:00+00:00 ---
Requesting window: 2026-08-10 13:30:00+00:00 -> 2026-08-10 14:30:00+00:00

Main dataframe row count after merge: 91

30-min moving averages updated.

=== Streaming simulation complete ===
```

Fig 5: The main pull loop — antijoin merge and moving average calculation.

4. Results

After all 7 pulls finished, here's what the two dataframes looked like.

```
... Final streamed dataframe (main_df):
```

UTC Timestamp	Stock Name	c	l	h	o	v
2026-08-10 13:00:00	AAPL	309.18	309.05	309.1986	309.05	11103.031054
2026-08-10 13:01:00	AAPL	309.2098	309.05	309.2098	309.07	5034.143329
2026-08-10 13:02:00	AAPL	308.94	308.94	309.19	309.1	16814.147142
2026-08-10 13:03:00	AAPL	308.96	308.9	309.06	308.9	4114.479768
2026-08-10 13:04:00	AAPL	309.0	308.91	309.0	308.94	3822.0
2026-08-10 13:05:00	AAPL	309.05	309.0	309.05	309.0	2072.31935
2026-08-10 13:06:00	AAPL	308.96	308.94	309.07	309.02	7473.551864
2026-08-10 13:07:00	AAPL	309.0	308.9	309.05	308.9252	4563.316874
2026-08-10 13:08:00	AAPL	309.0	308.97	309.1503	309.1499	2543.862054
2026-08-10 13:09:00	AAPL	308.91	308.9	309.01	309.01	1639.540547
2026-08-10 13:10:00	AAPL	309.04	308.9	309.04	308.9	2349.053553
2026-08-10 13:11:00	AAPL	309.25	309.0279	309.25	309.05	6286.504372
2026-08-10 13:12:00	AAPL	309.1373	309.0	309.2	309.1614	4146.395548
2026-08-10 13:13:00	AAPL	309.0977	309.07	309.1431	309.1431	1162.659916
2026-08-10 13:14:00	AAPL	309.05	309.001	309.12	309.09	5925.0
2026-08-10 13:15:00	AAPL	309.11	309.085	309.3	309.1	8796.57
2026-08-10 13:16:00	AAPL	309.19	309.1139	309.19	309.1762	1863.016171
2026-08-10 13:17:00	AAPL	309.2	309.15	309.2	309.1816	2403.876556
2026-08-10 13:18:00	AAPL	309.13	309.12	309.2	309.12	1809.336675
2026-08-10 13:19:00	AAPL	309.1662	309.11	309.2	309.11	1290.541491
2026-08-10 13:20:00	AAPL	309.11	309.1	309.19	309.17	2848.199909
2026-08-10 13:21:00	AAPL	308.9	308.89	309.12	309.1	12204.595755

```
...
```

Fig 6: Final streamed dataframe — 91 one-minute rows covering the replay window.

```
Final moving averages dataframe (ma_df):
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
[Stage 955:===== > (28 + 4) / 32]
```

Datetime	c_MA	l_MA	h_MA	o_MA	v_MA
2026-08-10 14:00:00	306.2433	305.8976	306.57126	306.25385339999997	221415.06695810007
2026-08-10 14:05:00	306.4509300000001	306.1731	306.6772166666667	306.4171867333333	176531.94647126665
2026-08-10 14:10:00	306.5968033333334	306.37422000000004	306.8162166666668	306.5822166666667	161265.8644214667
2026-08-10 14:15:00	306.63138	306.4222200000001	306.84039333333334	306.6180166666667	144124.55847139997
2026-08-10 14:20:00	306.69006666666667	306.51622333333336	306.87672333333325	306.67735000000005	127429.49182433335
2026-08-10 14:25:00	306.79007	306.62457333333333	306.97546	306.7899100000001	118480.42675013335
2026-08-10 14:30:00	306.78280333333333	306.6296933333333	306.9427933333333	306.77297666666675	107063.11025166667

Fig 7: Final moving-average dataframe — one row per pull once 30 minutes of data was available.

5. Workflow Summary

To summarize: the notebook pulls AAPL 1-minute bars every 5 minutes for 30 minutes (7 pulls total), and each pull gets merged into a running Spark dataframe using a left-anti-join antijoin, so any timestamps that overlap between pulls get refreshed instead of duplicated. After every merge, I compute a 30-minute trailing moving average of close/low/high/open/volume and append it to a separate dataframe, keyed by the end time of that window (the Datetime column). By the end, the streamed dataframe (main_df) has 91 rows covering 13:00–14:30 UTC on the replay day, and the moving-average dataframe (ma_df) has one row for every pull once there was at least 30 minutes of data to average over.

6. Code File

The notebook file aapl_streaming_analytics.ipynb is attached alongside this report.

Link-

<https://drive.google.com/file/d/1IzCetoh5yuHm5NZjMjXlxZaxa4szldnt/view?usp=sharing>

Part 2

Airflow Data Pipelining

Deliverable 2 of 2 — screenshots, answers, and code for Task 1, Task 2.1, and Task 2.2

Contents

- Task 1: Helloworld (Q1.1, Q1.2(1), Q1.2(2))
- Task 2.1: 12-Task DAG
- Task 2.2: Stock Price Fetching, Prediction, and Storage
- Closing Notes
- Part 2 Submission Checklist

Link to the Scripts-

https://drive.google.com/drive/folders/19CTvIMAKYGc2vJP_tGFRd9xRM-Pu34HY?usp=drive_link

Overview

This part of the document covers Part 2 of Assignment 2 - building and scheduling data pipelines with Apache Airflow, starting from a hello-world DAG and building up to a small machine learning pipeline that predicts stock prices.

I installed Airflow locally rather than on a GCP VM, since it's free and faster to iterate on. One thing worth mentioning up front: the handout was written for an older version of Airflow, so a couple of things had to be adjusted:

- Airflow's "Tree" view was replaced by the "Grid" view starting in version 2.3. Wherever the assignment asks for a Tree view screenshot, I took a Grid view screenshot instead, since it shows the same run history at a glance.
- Airflow 3.x removed the SequentialExecutor entirely and renamed the webserver command. Since Task 1 specifically requires SequentialExecutor, I installed Airflow 2.10.5 instead of the newest version.

I set up a Python 3.11 virtual environment (Airflow 2.x doesn't support the newer Python version my system had by default), installed apache-airflow==2.10.5, and confirmed executor = SequentialExecutor in the config file before starting.

Task 1: Helloworld

Q1.1 : Installing Airflow and Starting the Webserver/Scheduler

I ran the webserver and scheduler in two separate terminals, and used a third terminal for actual work commands, since Airflow's webserver and scheduler are both long-running processes that need to stay open.




Fig 1: Webserver in browser after a successful start.

```
PS C:\WINDOWS\system32> wsl -d ubuntu
tuf78@Sabbir:/mnt/c/WINDOWS/system32$ source ~/airflow_project/airflow_venv/bin/activate && export AIRFLOW_HOME=~/.airflow
w
(airflow_venv) tuf78@Sabbir:/mnt/c/WINDOWS/system32$ airflow scheduler

-----
  ____      _
 / ___|    / \
| |  | |  / _ \
| |  | | / ___ \
| |  | |/_/   \_\
| |  | |
|_|  |_|

[2026-08-14T02:45:49.410+0600] {executor_loader.py:258} INFO - Loaded executor: SequentialExecutor
[2026-08-14 02:45:49 +0600] [849] [INFO] Starting gunicorn 23.0.0
[2026-08-14 02:45:49 +0600] [849] [INFO] Listening at: http://[::]:8793 (849)
[2026-08-14 02:45:49 +0600] [849] [INFO] Using worker: sync
[2026-08-14 02:45:49 +0600] [850] [INFO] Booting worker with pid: 850
[2026-08-14T02:45:49.499+0600] {scheduler_job_runner.py:950} INFO - Starting the scheduler
[2026-08-14T02:45:49.500+0600] {scheduler_job_runner.py:957} INFO - Processing each file at most -1 times
[2026-08-14T02:45:49.513+0600] {manager.py:174} INFO - Launched DagFileProcessorManager with pid: 851
[2026-08-14T02:45:49.516+0600] {scheduler_job_runner.py:1949} INFO - Adopting or resetting orphaned tasks for active dag runs
[2026-08-14T02:45:49.518+0600] {settings.py:63} INFO - Configured default timezone UTC
[2026-08-14T02:45:49.536+0600] {scheduler_job_runner.py:1972} INFO - Marked 1 SchedulerJob instances as failed
[2026-08-14T02:45:49.554+0600] {manager.py:406} WARNING - Because we cannot use more than 1 thread (parsing_processes = 2) when using sqlite. So we set parallelism to 1.
[2026-08-14 02:45:49 +0600] [852] [INFO] Booting worker with pid: 852
[2026-08-14T02:45:50.046+0600] {dag.py:4180} INFO - Setting next_dagrun for task2_2_dag to 2026-08-13 00:00:00+00:00, run_after=2026-08-14 00:00:00+00:00
[2026-08-14T02:45:50.057+0600] {dag.py:4180} INFO - Setting next_dagrun for task2_1_dag to 2026-08-13 20:00:00+00:00, run_after=2026-08-14 00:00:00+00:00
```



[DAGs](#)
[Cluster Activity](#)
[Datasets](#)
[Security](#)
[Browse](#)
[Admin](#)
[Docs](#)
20:48 UTC

Do not use **SQLite** as metadata DB in production – it should only be used for dev/testing. We recommend using Postgres or MySQL. [View details](#) for more information.

Do not use the **SequentialExecutor** in production. [Click here](#) for more information.

DAGs

All 12
Active 3
Paused 69

Running 0
Failed 0

Auto-refresh

DAG ↕	Owner ↕	Runs ↕	Schedule	Last Run ↕	Next Run ↕	Recent Tasks ↕	Actions	Links
 Params Trigger UI example params	airflow	  	None			  	  	...
 Params UI tutorial example params ui	airflow	  	None			  	  	...
 Sample DAG with Display Name example	airflow	  	None			  	  	...
 conditional_dataset_and_time_based_timetable dataset time based timetable	airflow	  	Dataset or @ 11:3		2026-06-12, 01:00:00	  	  	...
 consume_1_and_2_with_dataset_expressions	airflow	  	Dataset		0 of 2 datasets updated	  	  	...
 consume_1_or_2_with_dataset_expressions	airflow	  	Dataset		0 of 2 datasets updated	  	  	...
 consume_1_or_both_2_and_3_with_dataset_expressions	airflow	  	Dataset		0 of 3 datasets updated	  	  	...
 dataset_alias_example_alias_consumer consumer dataset alias	airflow	  	Unresolved DatasetAlias			  	  	...
 dataset_alias_example_alias_consumer_with_no_taskflow consumer dataset alias	airflow	  	Unresolved DatasetAlias			  	  	...
 dataset_alias_example_alias_producer	airflow	  	None			  	  	...

Q1.2(1): Running Helloworld with SequentialExecutor

Note: as mentioned above, Airflow's Tree view was replaced by Grid view in 2.3+, so the screenshot below is Grid view in its place, it shows the same run-history-at-a-glance information.




Fig 4: Grid view for the helloworld DAG (substituting for the removed Tree view).




Fig 5: Graph view for the helloworld DAG.




Fig 6: Gantt view for the helloworld DAG.

Q1.2(2): Two Other Airflow UI Features

Besides Tree/Grid, Graph, and Gantt, I looked at two more tabs in the Airflow UI: Logs and Code. Both of these turned out to be genuinely useful later on, not just for this question, I actually used the Logs tab to debug a real failure in Task 2.1 (more on that below).

Logs Tab

You get to the Logs tab by clicking on a task instance box (for example, the hello_task box for helloworld), then clicking "Log." It shows the complete stdout/stderr output from that specific run of the task, for helloworld's hello_task, that includes the literal "Hello world!" print statement, plus Airflow's own metadata like start time, end time, and exit status. This is the main tool for figuring out why a task failed: instead of guessing, you can read the exact error or traceback that caused it. I relied on this heavily later when a task in my Task 2.1 DAG failed, the Logs tab is what showed me the actual exception.




Fig 7: Logs tab showing hello_task's output.

Code Tab

The Code tab shows the exact Python source file that Airflow parsed and is currently scheduling from. This is useful for confirming that a saved change to a DAG file was actually picked up, if the Code tab doesn't show your latest edit, it usually means either the file wasn't saved to the dags folder correctly, or the scheduler just hasn't rescanned yet (it rescans on a timer, not instantly), which was a source of confusion for me a couple of times while iterating on DAGs.




Fig 8: Code tab showing the helloworld DAG's source.

Logic:

- The DAG is defined using the `with DAG(...)` as dag: context manager (the modern Airflow style, vs. passing `dag=dag` to each operator explicitly), any operator instantiated inside this block is automatically registered to this DAG.
- `schedule_interval=None` means this DAG only runs when triggered manually, it won't fire on any automatic schedule, which fits Q1.2's requirement of running helloworld via manual execution first.
- `catchup=False` tells Airflow not to backfill any runs between `start_date` and today; since there's no schedule anyway, this is mostly a safe default here.
- The single task, `hello_task`, wraps `print_hello()` in a `PythonOperator`. The function both prints "Hello world!" to stdout (which lands in the task's log, visible in the Logs tab) and returns the same string, the return value gets automatically pushed to XCom, so even though nothing pulls it downstream here, it's a good example of Airflow's implicit XCom push behavior when writing about the Logs/XCom features later.
- With one task and no `>>` dependencies, the DAG graph is a single node, this makes it a clean smoke test to confirm the webserver, scheduler, and SequentialExecutor are all correctly wired before building the more complex 12-task DAG in Task 2.1.

Task 2.1: 12-Task DAG

Building the DAG

The assignment gives a diagram of a 12-task-looking DAG (it actually has 19 tasks, t1 through t19) with a specific dependency structure, and a color legend where yellow boxes are Python operators and gray boxes are Bash operators. I went through the diagram carefully and traced every arrow, since a few of the lines cross over each other and it's easy to misread which box connects to which.

For the operator types, the assignment asks for at least 3 different kinds of commands per operator type, so I did the following:

- Python tasks rotate between three functions: one that sleeps for a few seconds, one that just prints a message, and one that runs a small counting loop.
- Bash tasks rotate between three types of commands: a plain echo statement, calling a small shell script, and running a Python file.

I put the helper shell script and Python script in a separate `scripts/` folder next to the DAG file, so the DAG file itself stays focused on just defining the tasks and dependencies.

A Bug I Ran Into (and How I Fixed It)

The first time I triggered this DAG, one of the tasks failed almost immediately. Looking at the Logs tab for that task (this is exactly the Logs tab I described in Q1.2(2) above, it came in handy right away), I found a Jinja2 "TemplateNotFound" error, which was confusing at first since my bash command looked completely normal.

It turns out Airflow treats a BashOperator's command string as a Jinja2 template. If that command string happens to end in ".sh", Jinja's template loader misreads the entire command as the `FILENAME` of a template it should load from the `dags` folder, instead of just running it as a literal shell command. Since my command was something like "bash scripts/hello.sh", which ends exactly in ".sh", it tripped this behavior and Airflow tried (and failed) to find a template file with that whole string as its name.

The fix was: just made sure no `bash_command` string ends in ".sh". I added a trailing echo statement after every script call (for example, "bash scripts/hello.sh; echo done") so the command no longer ends in .sh, and after that change the task ran correctly.

(1) Tree/Graph Screenshots

Tree/Grid → Grid substitution note applies here as in Task 1.




Fig 9: Grid view for the 12-task DAG.




Fig 10: Graph view for the 12-task DAG, matching the assignment's diagram.

(2) Manual Trigger and Gantt

After fixing the bug above, I triggered the DAG manually and it completed successfully, all 19 tasks finished with no failures. Since I'm using SequentialExecutor, tasks run one at a time rather than in parallel, which you can see in the Gantt chart as non-overlapping bars.




Fig 11: Gantt chart after manually triggering the DAG, bars don't overlap since *SequentialExecutor* runs one task at a time.

(3) Scheduling

For the scheduling requirement, I set `start_date` to a fixed date in the past (`datetime(2026, 1, 1)`) with `catchup=False`, and `schedule_interval=timedelta(minutes=30)`.

The reasoning: because `start_date` is in the past but `catchup` is turned off, Airflow doesn't try to backfill every 30-minute interval that's already gone by, it only schedules the most recent interval. Since that interval boundary has already passed by the time I unpause the DAG, the first run fires immediately instead of me having to wait 30 minutes for it, and I also don't get flooded with historical catch-up runs. From there, `schedule_interval=timedelta(minutes=30)` just keeps firing a new run every 30 minutes.

I confirmed this worked as expected: after unpauseing, the DAG fired an immediate first run, then two more runs 30 minutes apart after that, three successful runs total, which satisfies "first run + 2 repeats and kept running for a day"

```
(airflow_venv) tuf78@Sabbir:/mnt/c/WINDOWS/system32$ airflow dags list-runs -d task2_1_dag
/home/tuf78/airflow_project/airflow_venv/lib/python3.11/site-packages/airflow/cli/commands/dag_command.py:48 UserWarning
: Could not import graphviz. Rendering graph to the graphical format will not be possible.
```

dag_id	run_id	state	execution_date	start_date	end_date
task2_1_dag	scheduled__2026-08-13 T20:30:00+00:00	success	2026-08-13T20:30:00+00:00	2026-08-13T21:00:01.3 95963+00:00	2026-08-13T21:01:05.982 999+00:00
task2_1_dag	scheduled__2026-08-13 T20:00:00+00:00	success	2026-08-13T20:00:00+00:00	2026-08-13T20:46:03.6 02027+00:00	2026-08-13T20:48:25.693 678+00:00
task2_1_dag	scheduled__2026-08-12 T22:00:00+00:00	success	2026-08-12T22:00:00+00:00	2026-08-13T20:45:50.0 78422+00:00	2026-08-13T20:48:19.553 640+00:00
task2_1_dag	scheduled__2026-08-12 T21:30:00+00:00	success	2026-08-12T21:30:00+00:00	2026-08-12T22:00:02.4 62722+00:00	2026-08-12T22:01:11.669 360+00:00
task2_1_dag	scheduled__2026-08-12 T21:00:00+00:00	success	2026-08-12T21:00:00+00:00	2026-08-12T21:30:01.5 22899+00:00	2026-08-12T21:31:10.528 098+00:00
task2_1_dag	scheduled__2026-08-12 T20:30:00+00:00	success	2026-08-12T20:30:00+00:00	2026-08-12T21:00:02.4 07231+00:00	2026-08-12T21:01:40.212 561+00:00
task2_1_dag	scheduled__2026-08-12 T20:00:00+00:00	success	2026-08-12T20:00:00+00:00	2026-08-12T20:30:01.5 06733+00:00	2026-08-12T20:31:02.269 414+00:00
task2_1_dag	manual__2026-08-12T19 :41:39.078600+00:00	success	2026-08-12T19:41:39.07 8600+00:00	2026-08-12T19:41:40.0 38408+00:00	2026-08-12T19:42:36.293 965+00:00
task2_1_dag	manual__2026-08-12T19 :34:53.389151+00:00	failed	2026-08-12T19:34:53.38 9151+00:00	2026-08-12T19:34:54.1 74096+00:00	2026-08-12T19:35:25.306 221+00:00
task2_1_dag	scheduled__2026-08-12 T19:30:00+00:00	success	2026-08-12T19:30:00+00:00	2026-08-12T20:00:02.3 01264+00:00	2026-08-12T20:01:01.456 850+00:00
task2_1_dag	scheduled__2026-08-12 T19:00:00+00:00	success	2026-08-12T19:00:00+00:00	2026-08-12T19:52:51.4 05138+00:00	2026-08-12T19:53:48.246 597+00:00

Fig 12: Run history showing the first scheduled run plus two repeats and more

Code-

```
1 import os
2 import time
3 from datetime import datetime, timedelta
4
5 from airflow import DAG
6 from airflow.operators.bash import BashOperator
7 from airflow.operators.python import PythonOperator
8
9 SCRIPTS_DIR = os.path.join(os.path.dirname(__file__), 'scripts')
10
11 default_args = {
12     'owner': 'sabbir',
13     'retries': 0,
14 }
15
16 def sleep_task(**kwargs):
17     tid = kwargs['task_instance'].task_id
18     print(f"Task {tid}: sleeping for 3 seconds")
19     time.sleep(3)
20     print(f"Task {tid}: done sleeping")
21
22 def print_task(**kwargs):
23     tid = kwargs['task_instance'].task_id
24     print(f"Task {tid}: Hello from print_task!")
25
26
```

```
25
26 def count_task(**kwargs):
27     tid = kwargs['task_instance'].task_id
28     count = 0
29     for i in range(10):
30         count += 1
31     print(f"Task {tid}: counted up to {count}")
32
33 with DAG(
34     dag_id='task2_1_dag',
35     default_args=default_args,
36     description='12-task DAG for Assignment 2 Task 2.1',
37     start_date=datetime(2026, 1, 1),
38     schedule_interval=timedelta(minutes=30),
39     catchup=False,
40     tags=['assignment2', 'task2.1'],
41 ) as dag:
42
43     t1 = BashOperator(task_id='t1', bash_command='echo "t1: starting the pipeline"')
44     t2 = BashOperator(task_id='t2', bash_command='echo "t2: running"')
45     t3 = BashOperator(task_id='t3', bash_command=f'bash {SCRIPTS_DIR}/hello.sh; echo "t3: script completed"')
46     t4 = PythonOperator(task_id='t4', python_callable=print_task)
47     t5 = BashOperator(task_id='t5', bash_command=f'python3 {SCRIPTS_DIR}/counter.py')
48     t6 = BashOperator(task_id='t6', bash_command='echo "t6: leaf task done"')
49     t7 = PythonOperator(task_id='t7', python_callable=sleep_task)
50
```

```
50     t8 = PythonOperator(task_id='t8', python_callable=count_task)
51     t9 = BashOperator(task_id='t9', bash_command='echo "t9: running"')
52     t10 = BashOperator(task_id='t10', bash_command=f'bash {SCRIPTS_DIR}/hello.sh; echo "t10: script completed"')
53     t11 = PythonOperator(task_id='t11', python_callable=print_task)
54     t12 = PythonOperator(task_id='t12', python_callable=sleep_task)
55     t13 = BashOperator(task_id='t13', bash_command=f'python3 {SCRIPTS_DIR}/counter.py')
56     t14 = BashOperator(task_id='t14', bash_command='echo "t14: merging branches"')
57     t15 = PythonOperator(task_id='t15', python_callable=count_task)
58     t16 = PythonOperator(task_id='t16', python_callable=print_task)
59     t17 = BashOperator(task_id='t17', bash_command=f'bash {SCRIPTS_DIR}/hello.sh; echo "t17: script completed"')
60     t18 = BashOperator(task_id='t18', bash_command=f'python3 {SCRIPTS_DIR}/counter.py')
61     t19 = PythonOperator(task_id='t19', python_callable=sleep_task)
62
63     t1 >> [t2, t3, t4, t5]
64     t2 >> t6
65     t3 >> t7
66     t5 >> [t8, t9]
67     t7 >> [t12, t13, t18]
68     t8 >> t10
69     t9 >> [t11, t12]
70     t10 >> t14
71     t11 >> [t14, t15]
72     t12 >> t14
73     t13 >> t18
74     t14 >> [t16, t17]
75
```

```
75     t15 >> t18
76     t16 >> t19
77     t17 >> t18
78     t18 >> t19
79
```

Here:

The DAG defines 19 tasks split between two operator types, BashOperator (11 tasks: t1, t2, t3, t5, t6, t9, t10, t13, t14, t17, t18) and PythonOperator (8 tasks: t4, t7, t8, t11, t12, t15, t16, t19), satisfying the "at least 3 different commands per operator type" requirement:

Bash tasks use three distinct command styles: plain echo statements (t1, t2, t6, t9, t14), calling an external shell script (hello.sh - t3, t10, t17), and running an external Python file via python3 (counter.py - t5, t13, t18).

Python tasks call one of three functions: sleep_task (simulates a delay via time.sleep(3)), print_task (logs a simple message), and count_task (runs a small loop and logs the result) - used across t4/t7/t8/t11/t12/t15/t16/t19.


Each Python callable accepts ****kwargs** and pulls task_instance from the Airflow context to log its own task_id, so the same three functions are reused across multiple tasks while still producing distinguishable log output per task.

The dependency graph is declared entirely with **>>** operators at the bottom, matching the assignment's diagram edge-for-edge (e.g., t1 **>>** [t2, t3, t4, t5] fans out from the root, t14 **>>** [t16, t17] merges branches back together before converging on t19 at the end). **schedule_interval=timedelta(minutes=30)** with **catchup=False** and **start_date=datetime(2026,1,1)** (a past date) makes the DAG fire immediately on unpause and then every 30 minutes, without backfilling historical runs.

Task 2.2: Stock Price Fetching, Prediction, and Storage

(1)–(2) Data Fetching and Preprocessing

I used yfinance's `Ticker(ticker).history(period='max')` to pull the full available price history for each of the 5 companies (AAPL, GOOGL, META, MSFT, AMZN), and saved each one to its own CSV file. For preprocessing, I kept only the columns I actually need (Open, High, Low, Close, Volume, and the Date), sorted everything by date, and dropped the Dividends and Stock Splits columns that yfinance adds but that the assignment doesn't ask for.

The image shows a code editor window titled 'fetch_data.py 1 X'. The background is dark with light-colored text. The code is as follows:

```
1 import os
2 import yfinance as yf
3
4 COMPANIES = ['AAPL', 'GOOGL', 'META', 'MSFT', 'AMZN']
5 DATA_DIR = os.path.join(os.path.dirname(__file__), 'data')
6
7 def fetch_all():
8     os.makedirs(DATA_DIR, exist_ok=True)
9     for ticker in COMPANIES:
10         df = yf.Ticker(ticker).history(period='max')
11         df = df.reset_index()
12         out_path = os.path.join(DATA_DIR, f'{ticker}.csv')
13         df.to_csv(out_path, index=False)
14         print(f"{ticker}: {len(df)} rows saved to {out_path}")
15
16 if __name__ == '__main__':
17     fetch_all()
18
```

Fig 14: Code for fetching and saving historical price data.

(3) Model Training- Walk-Forward Approach

For each company, I trained a linear regression model using [Open, High, Low, Close, Volume] on a given day as the input features, and the next day's High price as the target. The assignment asks for a walk-forward setup: for the last 5 trading days ($i = 5, 4, 3, 2, 1$), retrain the model using only data available up to $(d_{\text{current}} - i)$ days, then use that model to predict day $(d_{\text{current}} - (i-1))$'s high price. This gives 5 walk-forward predictions per company, and each prediction only ever uses a model trained on data that would actually have been available at that point in time, no lookahead.

```

train_predict.py X
Ubuntu > home > tuf78 > airflow > dags > scripts > task2_2 > train_predict.py > ...
1  import os
2  import pandas as pd
3  from sklearn.linear_model import Ridge
4
5  DATA_DIR = os.path.join(os.path.dirname(__file__), 'data')
6  FEATURES = ['Open', 'High', 'Low', 'Close', 'Volume']
7  RIDGE_ALPHA = 1.0 # small L2 regularization to counter OHLC multicollinearity
8
9  def load_company_df(ticker):
10     path = os.path.join(DATA_DIR, f'{ticker}.csv')
11     df = pd.read_csv(path)
12     df['Date'] = pd.to_datetime(df['Date'], utc=True).dt.date
13     df = df[['Date'] + FEATURES].sort_values('Date').reset_index(drop=True)
14     return df
15
16  def build_xy(df):
17     X = df[FEATURES].iloc[:-1].reset_index(drop=True)
18     y = df['High'].iloc[1:].reset_index(drop=True)
19     dates_for_y = df['Date'].iloc[1:].reset_index(drop=True)
20     return X, y, dates_for_y
21
22  def walk_forward_errors(ticker):
23     df = load_company_df(ticker)
24     X, y, dates_for_y = build_xy(df)
25
26     n = len(X)
27     errors = {}
28
29     for i in [5, 4, 3, 2, 1]:
30         train_end = n - i
31         test_idx = n - i
32         if train_end <= 0 or test_idx >= n:
33             print(f"{ticker}: not enough data for i={i}, skipping")
34             continue
35
36         X_train, y_train = X.iloc[:train_end], y.iloc[:train_end]
37         model = Ridge(alpha=RIDGE_ALPHA, solver='svd')
38         model.fit(X_train, y_train)
39
40         X_test = X.iloc[[test_idx]]
41         y_actual = y.iloc[test_idx]
42         y_pred = model.predict(X_test)[0]
43
44         rel_error = (y_pred - y_actual) / y_actual
45         pred_date = dates_for_y.iloc[test_idx]
46         errors[pred_date] = rel_error
47
48         print(f"{ticker} i={i}: trained on {train_end} rows, predicted {pred_date} -> pred={y_pred:.4f} actual={y_actual}")
49
50     return errors
51
52  if __name__ == '__main__':
53     for ticker in ['AAPL', 'GOOGL', 'META', 'MSFT', 'AMZN']:
54         walk_forward_errors(ticker)
55         print()
56

```

Fig 15: Code for the walk-forward training/prediction loop.

A Problem I Ran Into- and How I Diagnosed It

When I first ran this with a plain LinearRegression model, something looked very wrong. For META and MSFT, the relative errors were small and reasonable (around 1-2%), matching the kind of numbers shown in the assignment's example table. But for AAPL, GOOGL, and AMZN, the errors were huge, consistently around -77% to -85%, meaning the model's predictions were wildly off from the actual price.


Since this wasn't random noise but a very consistent, systematic error for 3 out of 5 companies, I checked the correlation between the Open/High/Low/Close columns for each company and also computed the condition number of the feature matrix (a measure of how numerically stable a regression on that data will be). All 5 companies came back with OHLC correlations above 0.999 and condition numbers between roughly 1.6×10^8 and 1.8×10^9 , meaning the feature matrix is nearly singular across the board, since Open/High/Low/Close are almost identical numbers on any given day.

In other words, this instability exists for all 5 companies, not just the 3 broken ones — it just happened to produce visibly bad coefficients for AAPL/GOOGL/AMZN and by chance produced reasonable-looking coefficients for META/MSFT. This is a classic multicollinearity problem, not a coding mistake.

To fix it, I switched from plain LinearRegression to Ridge regression (still a linear model, just with a small L2 penalty that stabilizes the coefficients when features are this collinear), using $\alpha=1.0$ and the 'svd' solver. After that change, all 5 companies' errors dropped to a sane ± 1 -2% range, consistent with each other and with the assignment's example table.

(4) Error Calculation

For each prediction, I calculated the relative error as $(\text{predicted_high} - \text{actual_high}) / \text{actual_high}$, and stored the results in a table with one row per prediction date and one column per company.



```

build_errors_csv.py 1 X
Ubuntu > home > tut78 > airflow > dags > scripts > task2_2 > build_errors_csv.py > ...
1 import os
2 import pandas as pd
3 from train_predict import walk_forward_errors
4
5 OUTPUT_DIR = os.path.join(os.path.dirname(__file__), 'data')
6 COMPANIES = ['AAPL', 'GOOGL', 'META', 'MSFT', 'AMZN']
7 COLUMN_NAMES = {
8     'AAPL': 'APPLE',
9     'GOOGL': 'GOOGLE',
10    'META': 'META',
11    'MSFT': 'MICROSOFT',
12    'AMZN': 'AMAZON',
13 }
14
15 def build_errors_csv():
16     all_errors = {} # ticker -> {date: error}
17     for ticker in COMPANIES:
18         all_errors[ticker] = walk_forward_errors(ticker)
19
20     # Union of all dates across tickers, sorted
21     all_dates = sorted(set(d for errs in all_errors.values() for d in errs.keys()))
22
23     rows = []
24     for d in all_dates:
25         row = {'Date': d}
26         for ticker in COMPANIES:
27             row[COLUMN_NAMES[ticker]] = all_errors[ticker].get(d, None)
28         rows.append(row)
29
30     df = pd.DataFrame(rows)
31     out_path = os.path.join(OUTPUT_DIR, 'errors.csv')
32     df.to_csv(out_path, index=True, index_label='')
33     print(f"Errors CSV written to {out_path}")
34     print(df)
35     return df
36
37 if __name__ == '__main__':
38     build_errors_csv()
39

```

Fig 16: Code that assembles the final errors table from all 5 companies.

	Date	APPLE	GOOGLE	META	MICROSOFT	AMAZON
0	8/7/2026	0.003158	0.009306	-0.00314	0.000567	-0.01027
1	8/10/2026	0.024748	0.001974	-0.01389	-0.01757	-0.0063
2	8/11/2026	0.001967	0.012392	-0.01552	0.014263	0.010653
3	8/12/2026	0.009245	0.006897	0.005419	0.01118	0.00858
4	8/13/2026	-0.00269	-0.0023	-0.01348	-0.00679	0.004061

Fig 17: Final errors CSV — relative error per company per day, for the last 5 trading days.

(5) The Airflow DAG, Cross-Task Communication, and Scheduler Setup

The pipeline is broken into 7 tasks total: one `fetch_data` task that pulls fresh data for all 5 companies, then 5 parallel `train_predict_<ticker>` tasks (one per company) that each run the walk-forward training/prediction described above, and finally one `build_errors_csv` task that combines everything into the final CSV.

Cross-Task Communication (XCom)

For passing data between tasks, I used Airflow's XCom mechanism. Each `train_predict_<ticker>` task pushes its own `{date: error}` results to XCom using `xcom_push()`, and the final `build_errors_csv` task pulls all 5 of those dictionaries back out using `ti.xcom_pull(task_ids=...)`, then assembles them into the final table. XCom is the standard way to pass small pieces of data between tasks in Airflow, it avoids relying on any shared global state, which wouldn't work reliably anyway once tasks might run on different workers.

Scheduler Setup

Since the assignment asks to schedule the fetching/prediction to run once a day, I set `schedule_interval='@daily'` with `start_date=datetime(2026, 1, 1)` and `catchup=False` on the DAG - the same `catchup=False` reasoning as Task 2.1 applies here: it stops Airflow from backfilling every daily interval since the start date, and instead just fires the most recent one immediately once the DAG is unpaused. I tested this in two stages: first, triggered the DAG manually to confirm the whole pipeline works correctly end-to-end (`fetch` → 5 parallel predictions → combined CSV), and only after that did I unpause it. Once unpaused, the scheduler picked it up and fired an automatic daily run on its own, on top of the earlier manual run, which confirmed the daily schedule is actually working and not just configured on paper.

```

1 import os
2 import sys
3 from datetime import datetime
4
5 from airflow import DAG
6 from airflow.operators.python import PythonOperator
7
8 SCRIPTS_DIR = os.path.join(os.path.dirname(__file__), 'scripts', 'task2_2')
9 sys.path.insert(0, SCRIPTS_DIR)
10
11 COMPANIES = ['AAPL', 'GOOGL', 'META', 'MSFT', 'AMZN']
12 COLUMN_NAMES = {
13     'AAPL': 'APPLE',
14     'GOOGL': 'GOOGLE',
15     'META': 'META',
16     'MSFT': 'MICROSOFT',
17     'AMZN': 'AMAZON',
18 }
19
20 default_args = {
21     'owner': 'sabbir',
22     'retries': 0,
23 }
24
25 def fetch_data_task(**kwargs):
26     from fetch_data import fetch_all
27     fetch_all()
28
29 def train_predict_task(ticker, **kwargs):
30     from train_predict import walk_forward_errors
31     errors = walk_forward_errors(ticker)
32     # XCom requires JSON-serializable values -> convert date keys to strings
33     errors_serializable = {str(d): float(e) for d, e in errors.items()}
34
35     kwargs['ti'].xcom_push(key=f'errors_{ticker}', value=errors_serializable)
36
37 def build_errors_csv_task(**kwargs):
38     import pandas as pd
39     ti = kwargs['ti']
40
41     all_errors = {}
42     for ticker in COMPANIES:
43         all_errors[ticker] = ti.xcom_pull(task_ids=f'train_predict_{ticker}', key=f'errors_{ticker}')
44
45     all_dates = sorted(set(d for errs in all_errors.values() for d in errs.keys()))
46
47     rows = []
48     for d in all_dates:
49         row = {'Date': d}
50         for ticker in COMPANIES:
51             row[COLUMN_NAMES[ticker]] = all_errors[ticker].get(d)
52         rows.append(row)
53
54     df = pd.DataFrame(rows)
55     out_path = os.path.join(SCRIPTS_DIR, 'data', 'errors.csv')
56     df.to_csv(out_path, index=True, index_label='')
57     print(f"Errors CSV written to {out_path}")
58     print(df)
59
60 with DAG(
61     dag_id='task2_2_dag',
62     default_args=default_args,
63     description='yfinance stock price fetching, prediction, and error tracking pipeline',
64     start_date=datetime(2026, 1, 1),
65     schedule_interval='@daily',
66     catchup=False,
67     tags=['assignment2', 'task2.2'],

```

```

67 ) as dag:
68
69     fetch = PythonOperator(
70         task_id='fetch_data',
71         python_callable=fetch_data_task,
72     )
73
74     predict_tasks = []
75     for ticker in COMPANIES:
76         t = PythonOperator(
77             task_id=f'train_predict_{ticker}',
78             python_callable=train_predict_task,
79             op_kwargs={'ticker': ticker},
80         )
81         predict_tasks.append(t)
82
83     build_csv = PythonOperator(
84         task_id='build_errors_csv',
85         python_callable=build_errors_csv_task,
86     )
87
88     fetch >> predict_tasks >> build_csv
89

```

Fig 18: Source code for the Task 2.2 DAG.




Fig 19: Graph view of the Task 2.2 pipeline.




Fig 20: Logs of the Task 2.2 pipeline.




Fig 21: Gantt of the Task 2.2 pipeline

Task Checklist:

- Task 1: webserver + scheduler terminal screenshots, browser DAGs list, Grid/Graph/Gantt for helloworld, 2 extra UI features (Logs, Code) with explanations.
- Task 2.1: Grid/Graph screenshots, Gantt after manual trigger, scheduling explanation + 3-run history screenshot, DAG source code.
- Task 2.2: code screenshots for all 4 files (fetch_data.py, train_predict.py, build_errors_csv.py, task2_2_dag.py), errors CSV, DAG graph screenshot, workflow/XCom/scheduler explanation.
- Every screenshot is legible and matches the claim next to it (e.g. the 3 scheduled runs are actually visible).